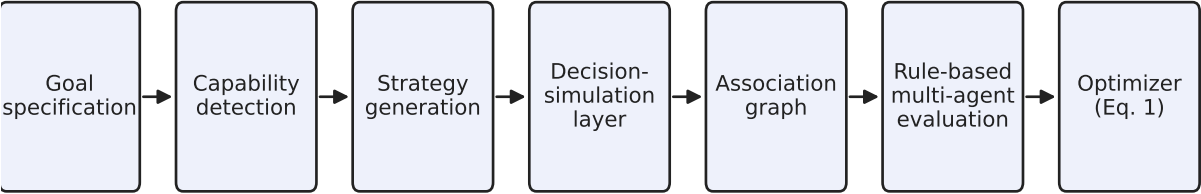


Figure 1. DecisionGPT pipeline and evaluation decomposition (schematic).



All stages are deterministic in every experiment reported here. When configured, an LLM parses the goal sentence and narrates the final result only — it does not compute agent scores or the selection in Eq. 1. No LLM was configured for any result in this paper. Explanation and memory run after selection and are not inputs to Eq. 1.