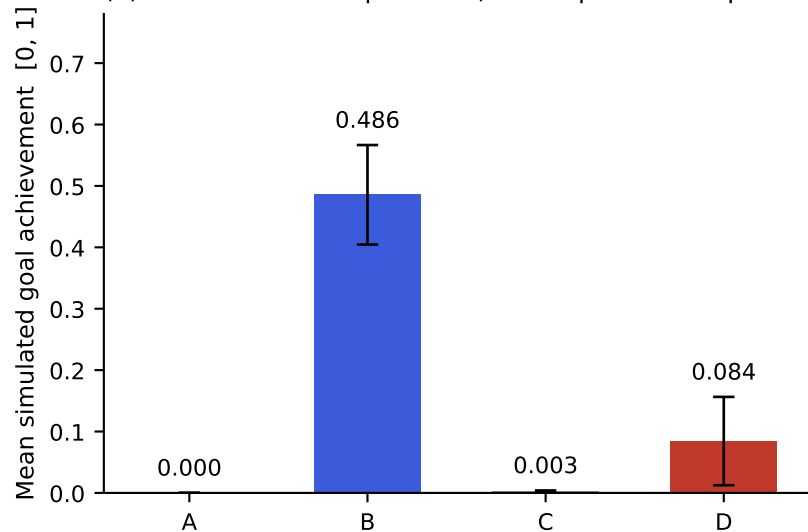
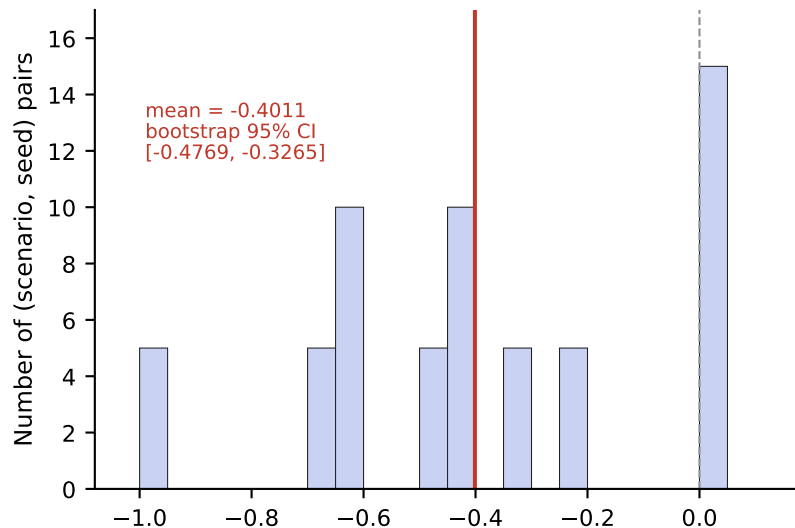


Figure 2. Primary result (experiment 0e1bd8dc): 12 designed synthetic scenarios $\times$ 5 seeds, paired by (scenario, seed). Simulated metric on procedurally generated businesses — not real-world effectiveness. Error bars in (a): Student-t 95% CI (within-suite variability). In (b) all 45 non-tied pairs favour B; Wilcoxon two-sided $p < 10^{-4}$ (recomputed 4.8×10^{-9}); rank-biserial = -1.00 .

(a) Architecture comparison (n = 60 paired obs. per bar)



(b) Paired differences, D vs B (0 / 15 / 45 D-better / tie / B-better)



A: prediction only B: + decision simulation C: + single agent D: full (+ multi-agent)

Paired difference D - B in simulated goal achievement